

# Mathematical Software Programming (02635)

Lecture 7 — October 23, 2025

Instructor: Martin S. Andersen

Fall 2025



# This week

## Topics

- ▶ timing your programs
- ▶ complexity
- ▶ basic computer architecture and efficiency
- ▶ compiler optimization

## Learning objectives

- ▶ analyze the runtime behavior and the time and space complexity of simple programs
- ▶ get a basic understanding of computer architecture and performance

# Timing your programs

Basic timing routines available in header file `time.h`

## Wall time

Prototype: `time_t time(time_t *tloc)`

- ▶ operating system time (not guaranteed to be *monotonic*)
- ▶ resolution: 1 second

## CPU time

Prototype: `clock_t clock(void)`

- ▶ returns (approximation to) processor time used by process
- ▶ resolution is system-dependent
- ▶ `clock()` may *wrap around*
- ▶ divide by macro `CLOCKS_PER_SEC` to convert to seconds

## Example: measuring CPU time

```
#include <time.h>

int main(void) {
    double cpu_time;
    clock_t T1, T2;
    T1 = clock();
    /* ... code you want to time ... */
    T2 = clock();
    cpu_time = ((double)(T2-T1)) / CLOCKS_PER_SEC;
    printf("CPU time: %.2e\n",cpu_time);
    return 0;
}
```

What can you do if time resolution is poor?

# Some platform-specific timing routines

## GNU/Linux (macOS 10.6+)

- ▶ `clock_gettime()` (`#include <time.h>`)

## POSIX (Portable Operating System Interface)

- ▶ `getrusage()` (`#include <sys/resource.h>`)

## macOS

- ▶ `mach_absolute_time()`, `mach_timebase_info()` (`#include <mach/mach_time.h>`)

## Windows

- ▶ `GetTickCount64()` (`#include <Windows.h>`)

# Profiling

Purpose: get the bigger picture

- ▶ Dynamic program analysis
- ▶ Find the hotspots/bottlenecks in your code
- ▶ Analyze space (memory) and time complexity of programs

## Tools

Modern (statistical) tools (no code change or re-compilation):

- ▶ Linux: Performance Tools (`perf`)
- ▶ macOS: Xcode / Instruments (`iprofiler`)
- ▶ Windows: Visual Studio Profiling Tools
- ▶ Google performance tools (`gperf-tools`)

Classical tools (instrumenting your code, needs re-compilation):


- ▶ Linux/Unix: `gprof` (see '`man gprof`' for more information)

# Big O notation

$f(x)$  is “big o” of  $g(x)$  if for some  $c > 0$  and  $x_0$

$$|f(x)| \leq c g(x), \quad \forall x \geq x_0,$$

written as  $f(x) = O(g(x))$  as  $x \rightarrow \infty$ .



$g(x)$  provides an **asymptotic upper bound** on  $f(x)$

## Big O notation (cont.)

Consider the function  $f(x) = 4x^3 - 2x^2 + 5x$  which satisfies

$$f(x) = 4x^3 - 2x^2 + 5x \leq |4x^3| + |2x^2| + |5x|.$$

Moreover, if  $x \geq 1$ ,

$$f(x) \leq |4x^3| + |2x^2| + |5x| \leq |4x^3| + |2x^3| + |5x^3| \leq 11x^3 \implies f(x) = O(x^3).$$

- ▶ the statements  $f(x) = O(x^4)$  and  $f(x) = O(2^x)$  are also true
- ▶ we are often interested in the “best” asymptotic upper bound
- ▶ “big omega” yields **asymptotic lower bound** on  $f$

$$f(x) = \Omega(g(x)) \iff g(x) = O(f(x))$$

- ▶ “big theta” yields **asymptotic tight bound** on  $f$

$$f(x) = \Theta(g(x)) \iff f(x) = O(g(x)) \wedge f(x) = \Omega(g(x))$$

# Time and space complexity

## Time complexity

- ▶ accessing  $i$ th element of an array requires  $O(1)$  time
- ▶ adding two vectors of length  $n$  requires  $O(n)$  time
- ▶ finding the maximum element of a vector requires  $O(n)$  time
- ▶ enumerating all possible permutations of  $n$  objects requires  $O(n!)$  time

## Space (memory) complexity

- ▶ a vector of length  $n$  requires  $O(n)$  memory
- ▶ a matrix of size  $m \times n$  requires  $O(mn)$  memory

## Some complexity classes (problem size $n$ )


Complexity	Big-O notation
Constant	$O(1)$
Logarithmic	$O(\log n)$
Poly-logarithmic	$O((\log n)^k)$
Linear	$O(n)$
Quasi-linear/log-linear	$O(n(\log n)^k)$
Polynomial	$O(n^k)$
Exponential	$O(2^{\text{poly}(n)})$
Factorial	$O(n!)$
Double exponential	$O(2^{2^{\text{poly}(n)}})$

## Comparison of time complexity

$T(n)$	$n = 10$	$n = 20$	$n = 30$	$n = 40$
$n$	0.00001 s	0.00002 s	0.00003 s	0.00004 s
$n^2$	0.0001 s	0.0004 s	0.0009 s	0.0016 s
$n^3$	0.001 s	0.008 s	0.027 s	0.064 s
$n^5$	0.1 s	3.2 s	24.3 s	1.7 min
$2^n$	0.001 s	1 s	17.9 min	12.7 days
$3^n$	0.059 s	58 min	6.5 years	3855 centuries

Hypothesis	Plot	Upper bound	Slope
$O(\log(n))$	Log $n$ -axis	$T(n) \leq c \log(n)$	$c$
$O(n^k)$	Log-log plot	$\log(T(n)) \leq k \log(n) + \log(c)$	$k$
$O(2^{\alpha n})$	Log $T$ -axis	$\log(T(n)) \leq \alpha \log(2) n + \log(c)$	$\alpha \log(2)$

# Running time $T(n)$ versus input size $n$



# Computer architecture


## Central processing unit

- ▶ CPU cores, math coprocessor
- ▶ registers and cache memory
- ▶ instruction pipelining

## Memory hierarchy

- ▶ L0: CPU registers
- ▶ L1: level 1 cache (SRAM)
- ▶ L2: level 2 cache (SRAM)
- ▶ L3: level 3 cache (SRAM)
- ▶ L4: random access memory (DRAM)
- ▶ L5: local secondary storage (local disks)
- ▶ L6: remote secondary storage (distributed file systems, servers)

## Computer architecture (cont.)



## Approximate latency comparison numbers (~2012)

Description	Time (ns)
CPU cycle	0.3
L1 cache reference	1
Branch mispredict	5
L2 cache reference	7
L3 cache reference	13
Main memory reference	100
Read 4K randomly from SSD	150,000
Read 1 MB sequentially from memory	250,000
Read 1 MB sequentially from SSD	1,000,000
Disk seek	10,000,000
Read 1 MB sequentially from disk	20,000,000
Send packet CA–Netherlands–CA	150,000,000

# What are my system specs?

## Windows Command Prompt

```
C:\>wmic cpu get L2CacheSize, L3CacheSize
```

Alternatively, download and use [CPU-Z](#) system profiler.

## macOS

```
$ system_profiler SPHardwareDataType  
$ sysctl -a | grep "hw.*cache"
```

## Linux/Unix

```
$ lscpu | grep cache  
$ cat /proc/cpuinfo  
$ cat /proc/meminfo
```

# Locality

## Temporal locality

- ▶ reuse of instructions/data within a short window of time
- ▶ instructions may be reused in a *tight* loop
- ▶ repeatedly using/referencing the same variables

## Spatial locality

- ▶ use of data elements within relatively close storage locations
- ▶ small stride access patterns (e.g. stride-1 access)
- ▶ execute instructions in sequence

## Example

```
sum = 0;  
for (i=0; i<n; i++)  
    sum += data[i];
```

## Data size and cache size effects on performance


```
extern double arr[]; // global array defined in another source file

int datasize1(int elem) {

    for (int i=0; i<elem; i++) arr[i] *= 3; // 1 FLOP


    /* return the number of memory accesses */
    return(elem);
}
```

## Data size and cache size (cont.)



Specs: Intel Xeon X5550 @ 2.67GHz, 32 kB L1, 256 kB L2, 8 MB L3 cache

## Data size and cache size (cont.)



Specs: Intel Xeon E5-2660 v3 @ 2.60GHz, 32 kB L1, 256 kB L2, 25 MB L3 cache

## Cache speed and spatial locality

```
extern int N;                                // length or array
extern double arr[];


int stridetest(int incr) {

    for (int i = 0; i < N; i += incr) arr[i] *= 3; // 1 FLOP

    /* return the number of memory accesses */
    return(N/incr);
}
```

with incr equal to 1,2,4,8, ...


## Cache speed and spatial locality (cont.)



- ▶ Optimized code: uses a switch-statement on `incr` to allow extra compiler optimizations
- ▶ Random access:  $\sim 108$  MFLOPS

# The memory mountain

Intel Core i5 @ 2.7 GHz, 32k L1, 256k L2, 3MB L3



## Instruction-level parallelism

```
#define N 67108864          // 64*1024*1024
int arr[2] = {1,1};

for (i=0; i<N; i++) {    // Loop 1
    arr[0]++;
    arr[1]++;
}

for (i=0; i<N; i++) {    // Loop 2
    arr[0] += arr[1];
    arr[1] += arr[0];    // S2 (must wait for S1)
}
```

# Optimizing compilers

## Enable code optimization

```
$ gcc source.c -Wall -Ox -o my_program
```

- ▶ -O0 — no optimization (default, best option for debugging)
- ▶ -O1 — most common forms of optimization
- ▶ -O2 — additional code optimization
- ▶ -O3 — most “expensive” code optimization (may increase size)
- ▶ -Os — code optimization that reduces size of executable
- ▶ -ffast-math — may violate IEEE standard

Makefile: add optimization flags to CFLAGS (e.g., CFLAGS=-Wall -std=c99 -O2)

## Compile for a specific/generic CPU type

- ▶ -march=native (use gcc -mcpu=help to see supported CPU types)
- ▶ Intel CPUs: -march=corei7, -march=x86-64-v4, -msse4.2, -mavx512f, ...

## Example: loop unrolling

```
for (i = 0; i < N; i++) y[i] = i;
```

can be re-written by compiler as

```
for(i = 0; i < (N - N%4); i+=4) {  
    y[i] = i;  
    y[i+1] = i+1;  
    y[i+2] = i+2;  
    y[i+3] = i+3;  
}  
/* clean-up loop */  
for(i = 4*(N/4); i < N; i++) y[i] = i;
```

- ▶ improves the “work to overhead” ratio
- ▶ GCC compiler flag `-funroll-loops`
- ▶ Clang: may be enabled with `-O1 (#pragma clang loop unroll_count(N))`

## Today's exercises, part I: timing your code

Reproduce some of the plots from this lecture

- ▶ Write a simple timing framework to measure time and performance.
- ▶ Apply to the function `datasize1()` from the lecture.
- ▶ Get the performance characteristics of your computer.
- ▶ Compare with the specifications.

## Today's exercises, part II: matrix-vector product

### Looking ahead

- ▶ Next module: external linear algebra libraries (BLAS and LAPACK)
- ▶ BLAS matrix-vector multiplication function: **dgemv()** (mnemonic name for “**d**ouble precision **ge**matrix **m**ultiplication”)
- ▶ After BLAS: **parallel** matrix-vector product with OpenMP (Bernd Dammann)
- ▶ Today's task: write your own versions of `dgemv()`, i.e., `my_dgemv1()`, ...

### Two-dimensional arrays and cache effects

- ▶ Matrices are “two-dimensional” but mapped to the one-dimensional memory address space
- ▶ This can lead to “good” or “bad” memory access with respect to performance
- ▶ Today's exercise: explore this for two different versions of the matrix-vector product

## Today's exercises II: matrix times vector (cont.)

### Computation

$$y \leftarrow \alpha Ax + \beta y, \quad A \in \mathbb{R}^{m \times n}, \quad x \in \mathbb{R}^n, \quad y \in \mathbb{R}^m$$

### Method 1 (row-oriented)

$$y_i \leftarrow \alpha \sum_{j=1}^n A_{ij}x_j + \beta y_i, \quad i = 1, \dots, m$$

### Method 2 (column-oriented)

$$y \leftarrow \beta y + \sum_{j=1}^n \alpha x_j A_{:,j}$$

where  $A_{:,j}$  is the  $j$ th column of  $A$