

# Mathematical Software Programming (02635)

Lecture 8 — October 30, 2025

Instructor: Martin S. Andersen

Fall 2025



# Announcement

## Assignment

- ▶ Posted on **DTU Learn**
- ▶ Due on November 19, 2025 (submit report on **DTU Learn** and code on **Autolab**)
- ▶ Please post your questions on **DTU Learn**

## Learning objectives

- ▶ call external (third party) programs and libraries
- ▶ design, implement, and document a program that solves a mathematical problem
- ▶ debug and **test** mathematical software

# This week

## Topics

- ▶ External libraries and basic linear algebra

## Learning objectives

- ▶ Call external (third party) programs and libraries

# Basic Linear Algebra Subroutines (BLAS)

- ▶ building blocks for linear algebra (*de facto* standard)
- ▶ started as a FORTRAN library (late 1970s)
- ▶ linear algebra *engine* in MATLAB, Python (Numpy, Scipy), R, Mathematica, ...
- ▶ high performance when optimized for a specific system
- ▶ many vendors: Intel MKL, AMD AOCL, Accelerate, OpenBLAS, ATLAS

# BLAS

## BLAS level 1 routines (1970s)

- ▶ vector operations, e.g.,

$$x^T y, \quad \|x\|_2, \quad x \leftarrow \alpha x, \quad y \leftarrow \alpha x + y$$

- ▶ use  $O(n)$  operations for vectors of length  $n$

## BLAS level 2 routines (1980s)

- ▶ matrix-vector operations, e.g.,

$$y \leftarrow \alpha A x + \beta y, \quad A \leftarrow \alpha x x^T + A, \quad x \leftarrow T^{-1} x, \quad T \text{ triangular}$$

- ▶ use  $O(mn)$  operations for matrices of size  $m \times n$

# BLAS

## BLAS level 3 routines (1980s)

- ▶ matrix-matrix routines, e.g.,

$$C \leftarrow \alpha AB + \beta C, \quad X \leftarrow T^{-1}X, \quad T \text{ triangular}$$

- ▶ use  $O(n^3)$  operations for matrices of size  $n \times n$

# What's in a name?

## BLAS naming scheme


XYYZZ

- ▶ First character X indicates data type (S, D, C, Z)
- ▶ BLAS level 1: letters YYZZ indicate mathematical operation
- ▶ BLAS level 2+3: letters YY indicate matrix type
- ▶ BLAS level 2+3: letters ZZ indicate mathematical operation


## Examples

- ▶ `dscal` — *double scale* ( $x \leftarrow \alpha x$ )
- ▶ `saxpy` — *single a x plus y* ( $y \leftarrow \alpha x + y$ )
- ▶ `dgemv` — *double general matrix-vector* ( $y \leftarrow \alpha Ax + \beta y$ )
- ▶ `dtrsv` — *double triangular solve vector* ( $x \leftarrow T^{-1}x$ )
- ▶ `ssymm` — *single symmetric matrix-matrix* ( $C \leftarrow \alpha SB + \beta C$ )


# BLAS performance (OpenBLAS)




## BLAS performance (Accelerate)




# BLAS performance (MKL)




# BLAS performance (OpenBLAS)




## BLAS performance (Accelerate)




## BLAS performance (OpenBLAS)



## BLAS performance (Accelerate)




# Memory hierarchy



## Possible bottlenecks


- ▶ CPU speed is a bottleneck — *CPU bound*
- ▶ cache size/speed is a bottleneck — *cache bound*
- ▶ memory size/speed is a bottleneck — *memory bound*
- ▶ disk or network speed is a bottleneck — *I/O bound*

## BLAS vector storage



## BLAS matrix storage: full

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
9	12	1			2	10	4			7	3	11			5	8	6		



GE (general), SY (symmetric), HE (Hermitian), TR (triangular)

## BLAS matrix storage: packed

0	1	2	3	4	5	6	7	8	9
9	2	1	5	10	4	8	3	6	7

	0	1	2	3
0	9			
1	2	10		
2	1	4	3	
3	5	8	6	7

TP (triangular packed), SP (symmetric packed), HP (Hermitian packed)

## BLAS matrix storage: band

$$A = \begin{bmatrix} a_{11} & a_{12} & & & & \\ a_{21} & a_{22} & a_{23} & & & \\ a_{31} & a_{32} & a_{33} & a_{34} & & \\ & a_{42} & a_{43} & a_{44} & a_{45} & \\ & & a_{53} & a_{54} & a_{55} & a_{56} \end{bmatrix}$$

5 × 6 matrix with  $l = 2$  subdiagonals  
and  $u = 1$  superdiagonal

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
	$a_{11}$	$a_{21}$	$a_{31}$	$a_{12}$	$a_{22}$	$a_{32}$	$a_{42}$	$a_{23}$	$a_{33}$	$a_{43}$	$a_{53}$	$a_{34}$	$a_{44}$	$a_{54}$		$a_{45}$	$a_{55}$			$a_{56}$			

0	1	2	3	4	5
	$a_{12}$	$a_{23}$	$a_{34}$	$a_{45}$	$a_{56}$
$a_{11}$	$a_{22}$	$a_{33}$	$a_{44}$	$a_{55}$	
$a_{21}$	$a_{32}$	$a_{43}$	$a_{54}$		
$a_{31}$	$a_{42}$	$a_{53}$			

## BLAS memory model

- ▶ vectors and matrices are contiguous arrays
- ▶ matrices are stored in column-major ordering
- ▶ *stride* refers to distance between *consecutive* elements
- ▶ *leading dimension* (LDA) refers to distance between columns

$$A = \begin{bmatrix} A_{11} & A_{12} & A_{13} & A_{14} & A_{15} \\ A_{21} & A_{22} & A_{23} & A_{24} & A_{25} \\ A_{31} & A_{32} & A_{33} & A_{34} & A_{35} \\ A_{41} & A_{42} & A_{43} & A_{44} & A_{45} \end{bmatrix}, \quad \begin{bmatrix} * & * & * & * & * \\ * & * & A_{23} & A_{24} & A_{25} \\ * & * & A_{33} & A_{34} & A_{35} \\ * & * & * & * & * \end{bmatrix}$$

- ▶  $i$ th column of  $A$  is a vector of length 4 with stride 1
- ▶  $i$ th row of  $A$  is a vector of length 5 with stride 4
- ▶  $(A_{11}, A_{22}, A_{33}, A_{44})$  is a vector of length 4 with stride 5
- ▶  $A$  is a matrix with 4 rows, 5 columns, stride 1, LDA 4
- ▶ *slice* (submatrix to the right) has 2 rows, 3 columns, stride 1, LDA 4
- ▶ see [Stride Visualizer](#)

## Example: BLAS level 1

```
void dscal_(
    const int * n,           /* Prototype for BLAS dscal */
    const double * a,         /* length of array      */
    double * x,              /* scalar a            */
    const int * incx          /* array x             */
);
int main(void) {
    int i,incx,n;
    double a, x[5] = {1.0,2.0,3.0,4.0,5.0};
    n = 5; a = 3.0; incx = 1;
    dscal_(&n, &a, x, &incx); // Scale the vector x by 3.0
    return 0;
}
```

# Makefile

Windows

```
LDLIBS=-lopenblas
```

macOS

```
LDLIBS=-lblas
```

DTU Unix system (Gbar)

```
LDFLAGS=-L/usr/lib64/atlas  
LDLIBS=-lsatlas
```

```
# Requires: "module load openblas"  
CPPFLAGS=-I$(MODULE_OPENBLAS_BASE_DIR)/include  
LDFLAGS=-L$(MODULE_OPENBLAS_BASE_DIR)/lib  
LDLIBS=-l$(MODULE_OPENBLAS_LIB)
```

## Example: CBLAS

```
#include <stdio.h>
#if defined(__MACH__) && defined(__APPLE__)
#include <Accelerate/Accelerate.h>
#else
#include <cblas.h>
#endif

int main(void) {
    int i,incx,n;
    double a, x[5] = {1.0,2.0,3.0,4.0,5.0};
    n = 5; a = 3.0; incx = 1;
    cblas_dscal(n, a, x, incx); // Scale the vector x by 3.0
    return 0;
}
```

# Makefile

Windows

```
LDLIBS=-lopenblas
```

macOS

```
LDLIBS=-lcblas
```

DTU Unix system (Gbar)

```
LDFLAGS=-L/usr/lib64/atlas  
LDLIBS=-lcblas -lsatlas
```

```
# Requires: "module load openblas"  
CPPFLAGS=-I$(MODULE_OPENBLAS_BASE_DIR)/include  
LDFLAGS=-L$(MODULE_OPENBLAS_BASE_DIR)/lib  
LDLIBS=-l$(MODULE_OPENBLAS_LIB)
```

# Linear Algebra Package (LAPACK)

Library of routines for numerical linear algebra (1990s)

- ▶ matrix factorizations (LU, Cholesky, QR, SVD, Schur, ...)
- ▶ systems of linear equations
- ▶ least-squares problems
- ▶ eigenvalue and singular value problems
- ▶ uses BLAS routines as building blocks

## Assignment

Self-contained program for solving a system of equations with LAPACK routine DGESV

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}, \quad b \in \mathbb{R}^n$$

# Today's exercises

## 1. Run BLAS/CBLAS examples

Linux and Windows WSL users: install OpenBLAS first

```
$ sudo apt-get install libopenblas-dev
```

## 2. Autolab

- ▶ BLAS level 1 (dscal, daxpy)
- ▶ BLAS level 2 (dtrsv)

Remember: BLAS assumes column-major storage!